

RAG_Local_HF_Weaviate_v3 (Draft)

Introduction & How to Use This Manual

Welcome to the Local RAG System Project

This manual walks you through creating a **Local Retrieval-Augmented Generation (RAG)** system using:

- **Weaviate** (vector database)
- **Hugging Face** (pre-trained AI models)
- **Python** (the glue that ties everything together)
- **Docker** (coming soon in Step 2)

A Companion to the TicTec Series

We're pairing this manual with the **TicTec** article series, which gives you context, storytelling, and daily how-to guides. The **manual** is your lab workbook—where you'll find the **exact steps** to follow along.

Quick Links to TicTec Articles (So Far)

- **Monday (Dec 23):** *When the Going Gets Weird, the Weird Turn Pro* — Big-picture intro to TicTec and the RAG project.
- **Tuesday (Dec 24):** *Don't Panic — Building Your First Local AI System* — Overview of the RAG architecture and your initial environment checklist.
- **Wednesday (Dec 25):** *Forks, Prefects, and GitHub Basics (Updated)* — Setting up GitHub, creating repositories, and version control fundamentals.

Note: We will reference these articles throughout the manual. For example, if you get stuck with GitHub basics, see **Wednesday's** piece for more detailed tips.

Step 1: Setting Up Your Environment

In this step, you'll ensure your local machine is ready to handle the code and tools needed for a RAG system. Before you spin up Docker or integrate Weaviate, you want to be 100% sure you've got the basics down: **Python**, **Git**, and **GitHub**.

1.1 Install & Verify Python

1. Download Python 3.9+

- Visit python.org/downloads and choose the latest stable version (3.9 or higher).
- On Windows, select “Add Python to PATH” during installation.
- Mac/Linux users typically follow their OS-specific installation instructions.

2. Verify Your Installation

Open a terminal or command prompt and type:

bash

Copy code

```
python --version
```

-
- You should see something like `Python 3.9.x`.

3. (Optional) Create a Test Script

If you’re new to Python, create a file called `hello.py` with:

python

Copy code

```
print("Hello, TicTec!")
```

-

Run it with:

bash

Copy code

```
python hello.py
```

-
- If it prints “Hello, TicTec!”, you’re golden.

TicTec Tie-In

- **Tuesday (Dec 24) article** introduced why Python is a critical part of building a local AI system and how you’ll use it to tie together retrieval and generation.
- Next week’s TicTec content will dig deeper into Python libraries, but for now, ensure Python runs smoothly.

1.2 Set Up Git & GitHub

GitHub is essential for version control and collaboration in this project—like your mission control center for all code and config files. If you followed **Wednesday’s TicTec article**, you’ve likely done this, but let’s restate the basics:

1. **Install Git** (if not already installed)
 - Visit git-scm.com and download for your OS.

Verify by typing:

bash

Copy code

```
git --version
```

- - You should see something like `git version 2.42.0`.
2. **Create a GitHub Account & Repository**
 - Go to [GitHub.com](https://github.com) → **Sign Up** (if you haven't already).
 - Create a **new repository** named `local_rag_project` (or similar).
 - Optionally initialize it with a README.
 3. **Clone or Initialize Locally**

Option A: Clone the Repo

bash

Copy code

```
git clone https://github.com/<YOUR-USERNAME>/local_rag_project.git
```

```
cd local_rag_project
```

-

Option B: Initialize a Folder Locally

bash

Copy code

```
mkdir local_rag_project
```

```
cd local_rag_project
```

```
git init
```

Then connect it to GitHub:

bash

Copy code

```
git remote add origin
```

```
https://github.com/<YOUR-USERNAME>/local_rag_project.git
```

```
git push -u origin main
```

-

4. **Make Your First Commit**

Add a simple file (e.g., `readme.md` or `hello.py`):

bash

Copy code

```
echo "TicTec RAG Project" > readme.md
```

```
git add .
```

```
git commit -m "Initial commit"
```

```
git push origin main
```

○

TicTec Tie-In

- **Wednesday (Dec 25):** *Forks, Prefects, and GitHub Basics* explains branching, pull requests, and collaboration. If you prefer a UI-based approach, that article walks you through it step-by-step.

1.3 Optional (But Recommended): Virtual Environments

Since you'll be installing Python libraries (like **transformers** and **weaviate-client**), a virtual environment keeps dependencies clean and contained.

Create a Virtual Environment

bash

Copy code

```
python -m venv venv
```

- 1.
2. **Activate It**
 - **Windows:** `venv\Scripts\activate`
 - **Mac/Linux:** `source venv/bin/activate`
3. **Document It**
 - Add a short line to your `readme.md` or project docs indicating how to activate the venv.
 - Consider adding a `.gitignore` entry for `venv` if you don't want to commit environment files to GitHub.

Why Bother?

- Avoid dependency conflicts (especially if you have multiple Python projects).

- Keep your RAG project self-contained, making it easier to replicate or move to another machine later.
-

Next Steps: A Teaser for Docker (Step 2)

You've laid the groundwork for your local RAG system. The next big piece is **Docker**, which helps containerize Weaviate (and potentially other services) so everything runs smoothly on your local machine—no more “works on my computer” problems.

- **TicTec will cover Docker** in upcoming articles (Week 2), walking you through installation, commands, and debugging.
- **This Manual** will detail Step 2 (Running Weaviate with Docker) after you're comfortable with the basics in the TicTec Docker week.

Sneak Peek

In Step 2, you'll:

- Pull the Weaviate Docker image
 - Run it locally on port 8080
 - Possibly experiment with `docker-compose.yml` to run multiple containers
-

Conclusion (Up to Docker)

Congratulations on completing **Step 1** of the *RAG_Local_HF_Weaviate_v3* manual! You have:

1. **Installed Python** and can run scripts locally.
2. **Set Up Git & GitHub** so you can track changes and collaborate.
3. (Optionally) **Created a Virtual Environment** for neat, conflict-free library installs.

When you're ready for more, we'll delve into **Docker** (Step 2) and how to containerize Weaviate. Until then, keep an eye on the next TicTec articles to reinforce these steps with real-world scenarios and hands-on GitHub workflows.

Quick Reference

- **Git Commands**
 - `git clone <URL>`
 - `git status`

- `git add .`
 - `git commit -m "Message"`
 - `git push origin main`
 - **Python Basics**
 - `python --version`
 - `python -m venv venv`
 - `pip install -r requirements.txt`
 - **Where to Learn More**
 - **TicTec Monday & Tuesday:** Big-picture overview of local RAG systems.
 - **TicTec Wednesday:** GitHub fundamentals, branching, and collaboration.
 - **TicTec Next:** Thursday & Friday deeper GitHub workflows, then into Docker next week.
-

End of “Up to Docker” Section

*(We'll update again once we're ready to integrate Docker specifics in **Step 2**.)*