

# RAG\_Local\_HF\_Weaviate\_v3 (TicTec Lab Workbook)

---

## Introduction & How to Use This Manual

### Welcome to the Local RAG System Project

This manual walks you through creating a **Local Retrieval-Augmented Generation (RAG)** system using:

- **Weaviate** (vector database)
- **Hugging Face** (pre-trained AI models)
- **Python** (the glue that ties everything together)
- **Docker** (coming soon in Step 2)

### A Companion to the TicTec Series

We're pairing this manual with the **TicTec** article series, which gives you context, storytelling, and daily how-to guides. The **manual** is your lab workbook—where you'll find the **exact steps** to follow along.

#### Quick Links to TicTec Articles (So Far)

- **Monday:** [When the Going Gets Weird, the Weird Turn Pro](#) — Big-picture intro to TicTec and the RAG project.
- **Tuesday:** [Don't Panic](#) — Why Python is critical and how it powers the RAG system.
- **Wednesday:** [Forks, Prefects, and GitHub Basics](#) — Setting up GitHub and mastering version control.
- **Thursday:** [Taming Thursdays – Diving Deeper into GitHub](#) — Cloning repositories and daily workflows.
- **Friday:** [Forks and Fridays – Taking the Path with GitHub Branches](#) — Mastering branching workflows for experimentation and collaboration.

**Note:** We will reference these articles throughout the manual. For example, if you get stuck with GitHub basics, see **Wednesday's** piece for more detailed tips.

## RAG\_Local\_HF\_Weaviate\_v3 (Draft)

---

# Step 1: Setting Up Your Environment

## Introduction: Building the Foundation

Welcome to your journey of creating a Local Retrieval-Augmented Generation (RAG) system! In this step, we'll set up your local development environment to handle the tools and code required for the project. By the end of this step, you'll have:

- A working installation of Python.
- Version control with Git and GitHub.
- (Optionally) a virtual environment to manage dependencies cleanly.

This step ensures you're ready to tackle containerization and other advanced concepts in future steps.

---

## 1.1 Install & Verify Python

### Why Python Matters

Python is the glue of your Local RAG system. It connects your vector database (Weaviate), pre-trained AI models (Hugging Face), and containerized services (Docker).

### Installation Steps

1. **Download Python 3.9 or Higher:**
  - Visit [python.org/downloads](https://python.org/downloads) and select the latest stable version (3.9 or higher).
  - During installation on Windows, check the box for "Add Python to PATH."
2. **Verify Your Installation:**
  - Open a terminal or command prompt and type:  
`python --version`
  - You should see output like `Python 3.9.x`.
3. **Test Python Installation:**

- Create a file named `hello.py` with the following content:  

```
print("Hello, TicTec!")
```
- Run the script:  

```
python hello.py
```
- If it prints "Hello, TicTec!", you're ready to move on.

### TicTec Tie-In

- **Article Link:** [Tuesday: Don't Panic](#)
  - *"Python is your key to connecting retrieval and generation in AI workflows."* This article explains why Python is critical for building local AI systems.
- 

## 1.2 Set Up Git & GitHub

### Why Version Control Is Essential

GitHub is your mission control for this project. It tracks code changes, facilitates collaboration, and keeps your files organized. Even for solo projects, version control is a lifesaver.

### Steps to Get Started

#### 1. Install Git:

- Download Git for your OS from [git-scm.com](https://git-scm.com).
- Verify installation:  

```
git --version
```
- You should see output like 

```
git version 2.42.0
```

.

#### 2. Create a GitHub Account:

- Go to [github.com](https://github.com) and sign up if you don't already have an account.

### 3. Create a Repository:

- Log in to GitHub and create a new repository named `local_rag_project` (or similar).
- Optionally initialize it with a README file.

### 4. Clone or Initialize Locally:

- **Option A:** Clone the repository:  

```
git clone  
https://github.com/<YOUR-USERNAME>/local_rag_project.git  
  
cd local_rag_project
```
- **Option B:** Initialize locally and connect to GitHub:  

```
mkdir local_rag_project  
  
cd local_rag_project  
  
git init  
  
git remote add origin  
  
https://github.com/<YOUR-USERNAME>/local_rag_project.git
```

### 5. Make Your First Commit:

Create a simple file (e.g., `readme.md`):

```
echo "TicTec RAG Project" > readme.md  
  
git add .  
  
git commit -m "Initial commit"  
  
git push origin main
```

## TicTec Tie-In

- **Article Link:** [Wednesday: Forks, Prefects, and GitHub Basics](#)
  - *"GitHub is your collaboration hub for the RAG project."* This article explains branching, pull requests, and using GitHub effectively.
- 

## 1.3 (Optional) Create a Virtual Environment

### Why Use a Virtual Environment?

Virtual environments keep your dependencies clean and isolated, especially when working on multiple Python projects. This ensures your RAG system remains self-contained and portable.

### Steps to Create and Activate a Virtual Environment

#### 1. Create the Environment:

```
python -m venv venv
```

#### 2. Activate It:

- **Windows:**  
`venv\Scripts\activate`
- **Mac/Linux:**  
`source venv/bin/activate`

#### 3. Document It:

- Add activation instructions to your `readme.md` or project documentation.
- Add `venv` to `.gitignore` to exclude it from version control:

```
venv/
```

## TicTec Tie-In

- **Article Link:** [Tuesday: Don't Panic](#)
  - *"Virtual environments keep your Python dependencies neat and conflict-free."* This article offers additional insights into maintaining clean development setups.
- 

## Quick Links to TicTec Articles

- **Monday:** [When the Going Gets Weird, the Weird Turn Pro](#) — Big-picture intro to TicTec and the RAG project.
  - **Tuesday:** [Don't Panic](#) — Why Python is critical and how it powers the RAG system.
  - **Wednesday:** [Forks, Prefects, and GitHub Basics](#) — Setting up GitHub and mastering version control.
  - **Thursday:** [Taming Thursdays – Diving Deeper into GitHub](#) — Cloning repositories and daily workflows.
  - **Friday:** [Forks and Fridays – Taking the Path with GitHub Branches](#) — Mastering branching workflows for experimentation and collaboration.
- 

## Next Steps: A Preview of Docker

Congratulations! You've completed Step 1 of your RAG system setup. Here's what you've accomplished:

- Installed Python and verified it works.
- Set up Git and GitHub for version control.
- (Optionally) created a virtual environment for clean dependency management.

In **Step 2**, we'll introduce Docker, the powerful tool that will containerize services like Weaviate and Elasticsearch. Keep following the TicTec series for context and detailed walkthroughs as we dive into containerization next.

---

## Docker section of RAG\_Local\_HF\_Weaviate\_v3 (TicTec Lab Workbook)

---

## Step 2: Introduction to Docker

## Picking Up Where We Left Off

Welcome back to your RAG system journey! In Step 1, you set up Python, Git, and GitHub, creating the foundation for your Local RAG system. Now it's time to tackle Docker—the tool that ensures your environment is as consistent as your code. With Docker, we'll eliminate the classic "works on my machine" problem and make running services like Weaviate and Elasticsearch seamless and predictable.

Before diving in, remember that this manual is designed to complement the **TicTec Docker Week articles**, which provide additional insights, storytelling, and practical examples. Refer to them whenever you'd like more context or a narrative perspective.

## 2.1 Why Docker Matters

### The Problem

Imagine sharing your Python application with a teammate or deploying it to a new server, only to encounter missing libraries, version mismatches, or OS-specific quirks. Sound familiar? These environment inconsistencies can grind development to a halt.

### The Solution

Docker solves this by packaging your application and its dependencies into portable, isolated containers. These containers run consistently, whether on your laptop, your teammate's machine, or in production.

### Quick Callout to TicTec

- **Article Link:** Monday: Why Docker Matters **(time travelers only)**
- *"Docker is the glue that holds development workflows together."* Use this article to understand how Docker brings order to chaotic environments.

---

## 2.2 Installing Docker

### Step-by-Step Installation

Docker installation depends on your operating system. Follow these steps:

#### For Windows:

1. **Download Docker Desktop:** Visit [docker.com](https://docker.com) and download Docker Desktop for Windows.
2. **Enable WSL 2:** Ensure Windows Subsystem for Linux 2 (WSL 2) is enabled. Follow [Microsoft's WSL 2 guide](#).

3. **Run the Installer:** Launch the Docker Desktop installer and follow the prompts.
4. **Verify Installation:** Open PowerShell and type:

```
docker --version
docker run hello-world
```

You should see Docker's "Hello, World!" message.

#### For macOS:

1. **Download Docker Desktop:** Visit [docker.com](https://docker.com) and download Docker Desktop for macOS.
2. **Install and Launch:** Drag Docker into your Applications folder, then open it.
3. **Verify Installation:** Open a terminal and type:

```
docker --version
docker run hello-world
```

#### For Linux:

1. **Install Docker Engine**

```
curl -o /tmp/get-docker.sh https://get.docker.com
sh /tmp/get-docker.sh
sudo usermod -aG docker $USER
```

2. **Verify Installation:**

```
docker --version
docker run hello-world
```

#### Troubleshooting Tips

- **Common Issues:** Port conflicts, firewall permissions, or virtualization disabled in BIOS.
- **Refer to TicTec:** Tuesday: Installing Docker Without Drama **(time travelers only)** for detailed troubleshooting steps.

---

## 2.3 Your First Docker Container



## Creating a Simple Dockerfile

Dockerfiles are the blueprint for containers. Let's create a simple one:

1. Create a directory and a file named `Dockerfile`.
2. Add this content:

```
FROM alpine:latest

CMD ["echo", "Hello, Docker World!"]
```

3. Build the image:

```
docker build -t hello-docker .
```

4. Run the container:

```
docker run hello-docker
```

You should see: `Hello, Docker World!`

## TicTec Tie-In

- **Article Link:** Wednesday: Building Your First Container **(time travelers only)**
- *"Building a container is like a magical first step into Docker."*
- This article dives deeper into container lifecycle management and explores key Docker concepts.

---

## 2.4 Writing Dockerfiles for Local RAG Systems

### Crafting an Effective Dockerfile

For our Local RAG project, we'll use a Dockerfile to containerize a Python app. Here's an example:

```
# Use a lightweight Python image
FROM python:3.9-slim
```

```
# Set the working directory
WORKDIR /app

# Copy dependency file and install libraries
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy the app code
COPY . .

# Set environment variables
ENV FLASK_APP=app.py

# Define the default command
CMD ["flask", "run", "--host=0.0.0.0"]
```

## Best Practices

- Minimize layers by combining commands.
- Use `.dockerignore` to exclude unnecessary files.
- Avoid hardcoding sensitive information (e.g., API keys).

## TicTec Tie-In

- **Article Link:** Thursday: Writing and Understanding Dockerfiles **(time travelers only)**
  - *"A well-crafted Dockerfile is your recipe for reliable containers."*
- 

## 2.5 Running Elasticsearch in Docker

### Getting Started with Elasticsearch

Elasticsearch is critical for RAG systems. Let's containerize it:

#### 1. Pull the Image:

```
docker pull docker.elastic.co/elasticsearch/elasticsearch:8.10.2
```

## 2. Run the Container:

```
docker run -d \  
  --name es-container \  
  -p 9200:9200 \  
  -e "discovery.type=single-node" \  
  docker.elastic.co/elasticsearch/elasticsearch:8.10.2
```

## 3. Verify:

- Logs: `docker logs es-container`
- Browser: Visit <http://localhost:9200>

## TicTec Tie-In

- **Article Link:** Friday: Running Elasticsearch in a Container **(time travelers only)**
  - *"Elasticsearch in Docker showcases the essence of portability and consistency."*
  - Check out the article for troubleshooting tips and real-world use cases.
- 

## Quick Links to TicTec Articles **(time travelers only)**

- **Monday:** Why Docker Matters
  - **Tuesday:** Installing Docker Without Drama
  - **Wednesday:** Building Your First Container
  - **Thursday:** Writing and Understanding Dockerfiles
  - **Friday:** Running Elasticsearch in a Container
- 

## Next Steps: Docker + Python Integration

You've completed Docker basics and are now ready to integrate containerized components with Python scripts in Step 3. Keep following the TicTec series for more insights and practical examples as we dive deeper into connecting Weaviate, Elasticsearch, and your local RAG system.

Congratulations on mastering Step 2—your RAG system is taking shape!

---